

Cross-dataset accuracy gap vs class-distribution skew
(binary Dir($\alpha = 0.1$); positive = KAN better)

